


LECTURE COMPUTER ARCHITECTURE


# ADVANCED MICROPROCESSOR ARCHITECTURES: INTEL X86

RAINER KELLER



## CONTENT

- 1 Overview of Microarchitectures
- 2 Internals of Intel
- 3 Internals of ARM

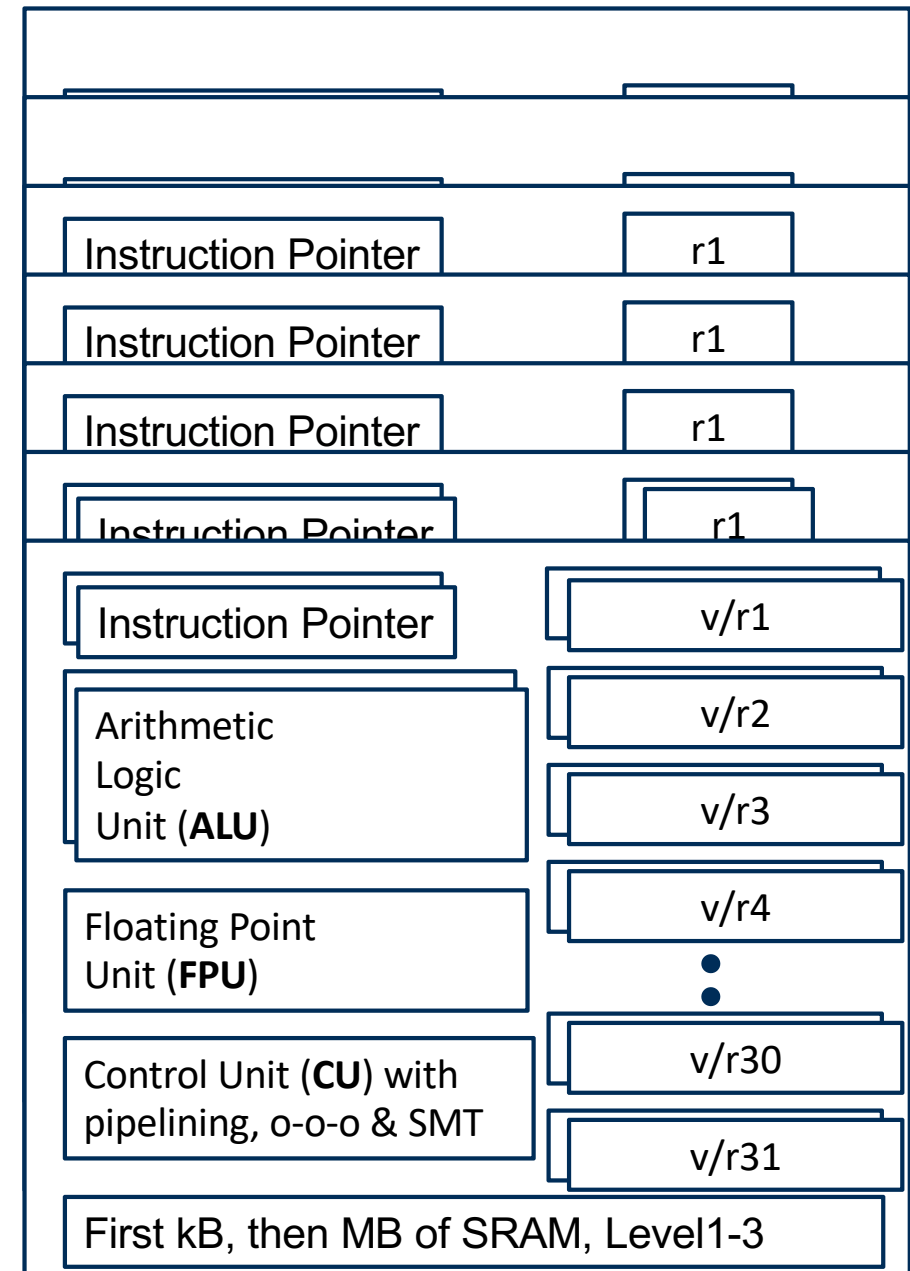


# Advanced Topics of CPUs

# REPETITION MICROPROCESSORS

Steps for improvement: **Check 01\_CA\_\***

1. Instruction Pipelining (shorter circuits → higher MHz, complex CU, from 5 → 34 stages)
2. Add integrated Cache (fast SRAM decouples DRAM, may be I-/D-Cache, multiple levels)
3. Multiple Instructions per Cycle (duplicate units in ALU, or complete ALU and FPU)
4. Speculative Out-of-Order Execution (compensate for bad instruction order and possibly add more register-files!)
5. Hyperthreading / SMT (not 1 thread of execution, but have HW track 2 threads; needs duplicated state, i.e. 2 IPs, two register files, etc., complex CU)
6. MMX, SSE, AVX (Single-Instruction-Multiple Data, SIMD): vector-operations, e.g. Intel AVX-512: 8 FP doubles, 16 FP floats per register...
7. multi-core.... **Many**-core (not shown here)



# Pipelining



# PIPELINING

- Pipelining visualized:  
Ford Model-T Assembly Line (1928):  
one worker repeats one assembly step.  
He's trained to be efficient & effective.  
All complete their step at ~ same time.



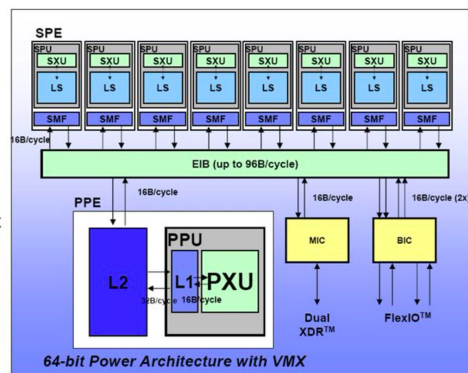
Source: alf-img.com

## IBM's Cell processor with 1x PowerPC and 8 "SPUs" in the Sony PS3:

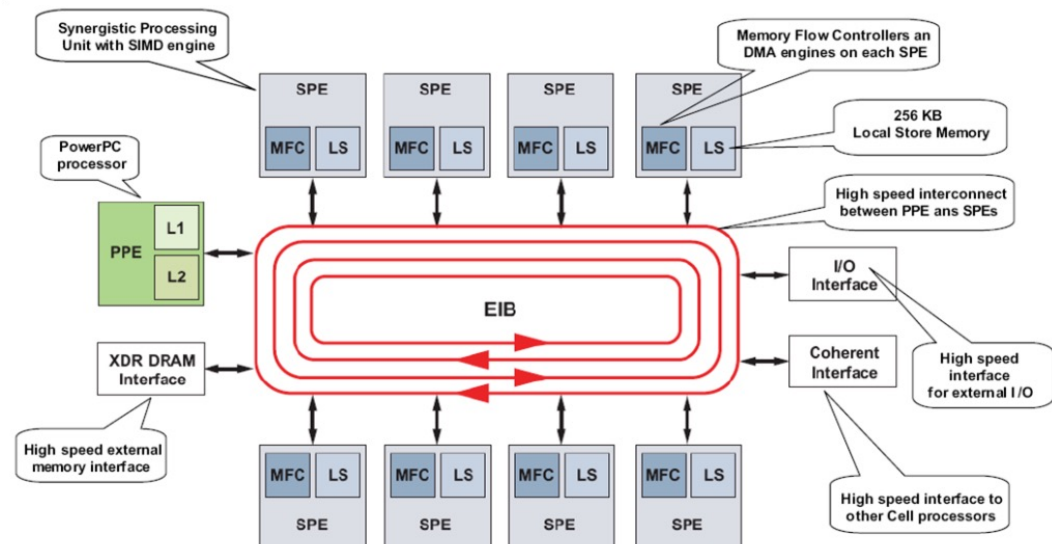
### IBM Cell Broadband Engine Architecture

- 3.2 GHz
- 9 cores, 10 threads
- >200 Gflops(single precision)
- >20 Gflops(double precision)
- Up to 25 GB/s memory bandwidth
- Up to 75 GB/s I/O bandwidth
- >300 GB/s interconnect bus

SPE: Synergistic Processor Element  
SPU: Synergistic Processor Unit  
SXU: SPU Core  
LS: Local Storage  
SMF: Synergistic Memory Flow Control  
EIB: Element Interconnect Bus  
PPE: PowerPC Processor Element  
PPU: Power Processor Unit  
PXU: Power Processor Unit  
MIC: Memory Interface Controller  
BIC: Bus Interface Controller  
L1: Memory Cache Internal to the CPU  
L2: Memory Cache External to the CPU



[http://domino.research.ibm.com/comm/research.nsf/pages/arch\\_innovation.htm](http://domino.research.ibm.com/comm/research.nsf/pages/arch_innovation.htm)



Source: redgamingtech.com / PS3 Post Mortem

Source: SlideShare, talk Michael A. Baker



## PARALLELISM VIA PIPELINING

Pipelining is an important concept in many areas – not just CS!

To improve instruction execution pipelining allows the overlap of many specialized units more efficiently, **even within the CPU!**

Instructions like `add r16, 0x20` may be divided into multiple smaller steps, i.e. these five phases, instead of sequential, pipeline them:

1. Instruction Fetch (IF)
2. Instruction Decode (ID)
3. Fetch Operands (OP)
4. Execute (EX)
5. Write Back (WB)



- This works only, if the instructions are **independent**; *the common case*
- If this is **not** the case, dependent instructions wait one or more cycles (so-called bubbles). In bad cases, the pipeline has to be **flushed**!

## PARALLELISM USING PIPELINING 2/2

If consecutive instructions are dependent instructions; CPU has to wait one (or more) cycles (so called “pipeline stalls/bubbles”):

```
add  r16,  0x20
mul  r16,  0x2
```

The second instruction uses the same register `r16` for addition, as the multiplication instruction, which therefore has to wait...


Even worse are „unexpected“ events, like

- function calls,
- (unpredicted) jumps
- hardware interrupts

leading to “pipeline flushes”, aka the whole pipeline needs to be rebuilt starting from the then new processor instruction.

Today the Branch-Prediction Unit tries to eliminate unnecessary stalls!  
(together with Out-of-Order Execution speculatively fetches & executes)

→ Leads to non-determinism...

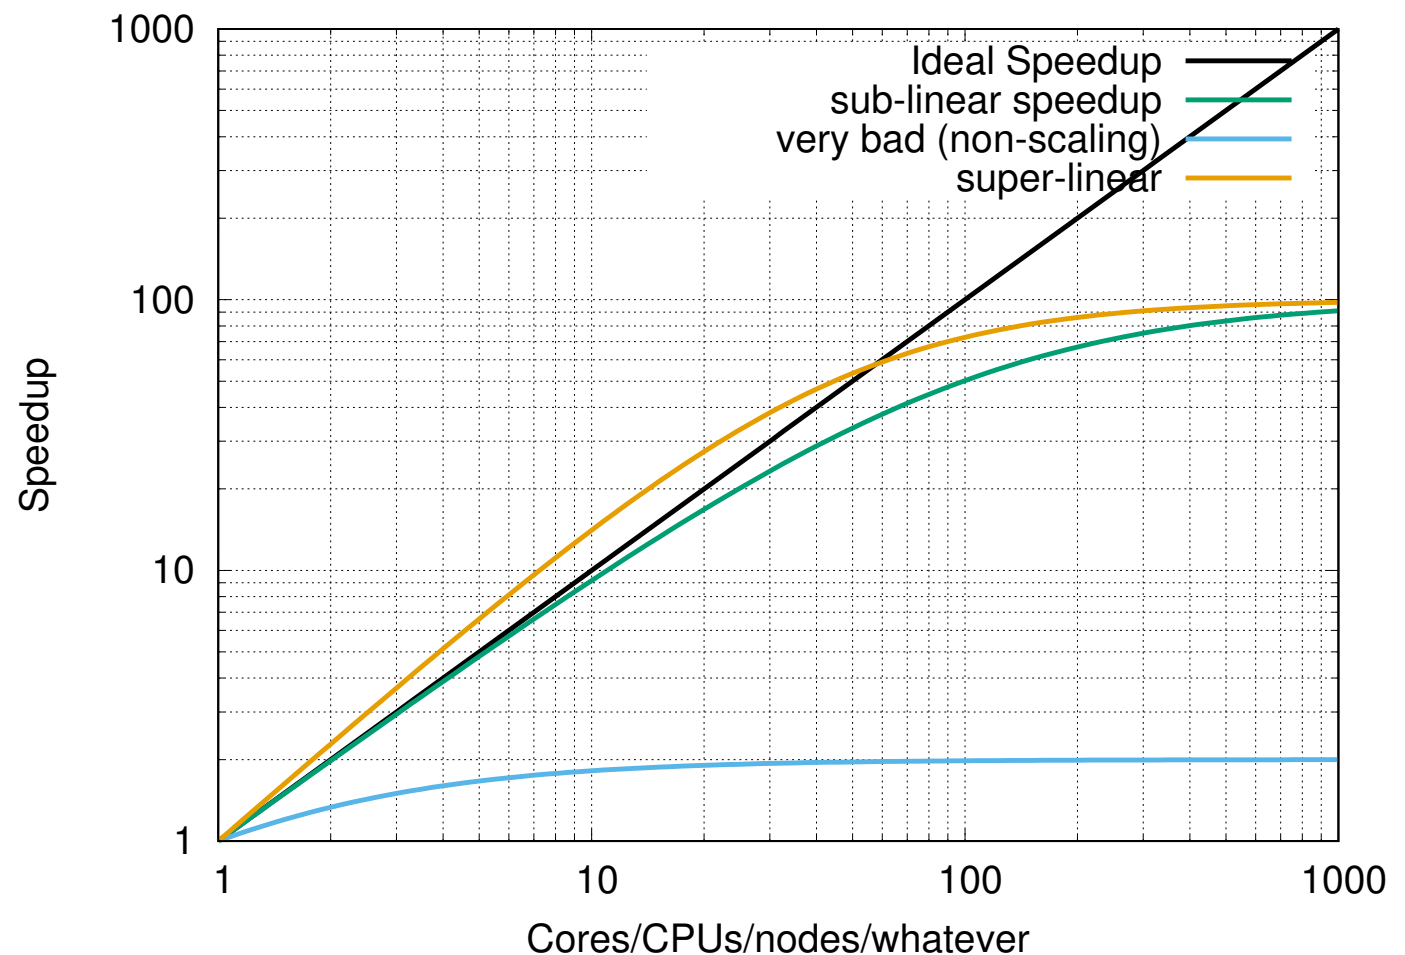




# Scaling

# SCALING

Building construction: digging 10m ditch takes 1 person ~10hrs  
How long does it take with 10 persons?



# Caches

# CACHES: SOLUTION TO AN EVER GROWING PROBLEM

A “cache” is a collection of data duplicating / shadowing data, usually to speedup the access, e.g. memory access.

A processor runs with over 3 GHz clock frequency; data has to be fetched from main memory at least once, i.e. DRAM

Access to (approximation):

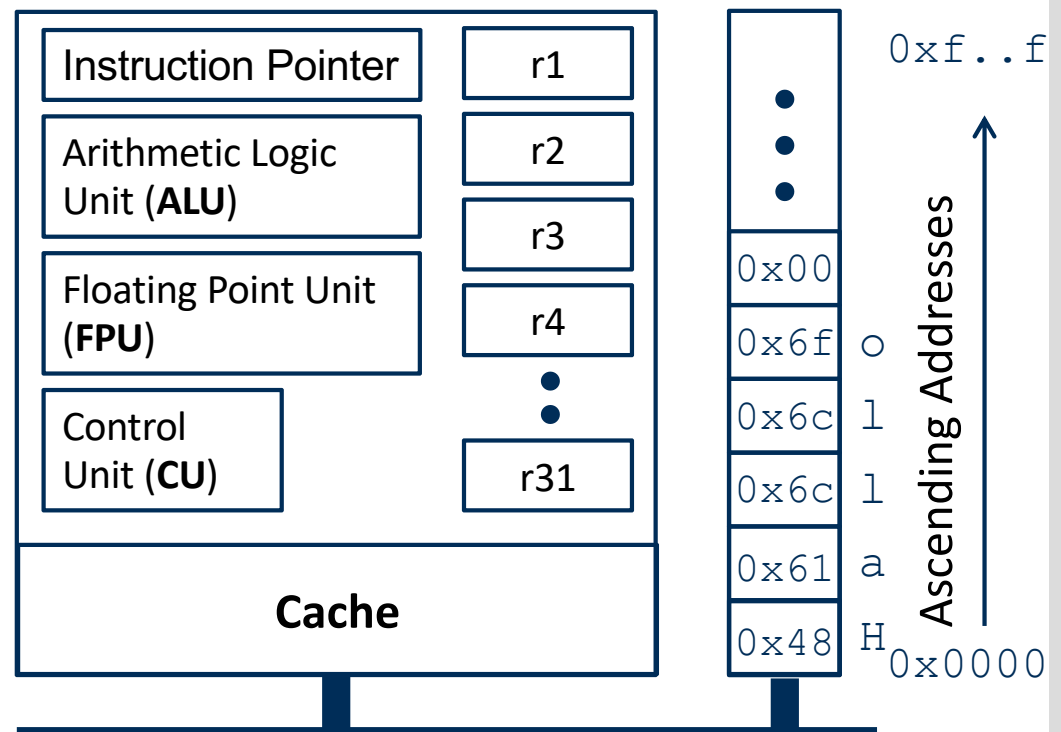
- Register: 1 cycle
- Cache: 1 - 10 cycles
- Main memory: >>100 cycles

To be fast, caches are:

- Small (1kB - few MB)
- Organized (see below)
- Tailored specifically to use-case.

Caches **only work**, if our code has:

- Temporal & Spatial locality



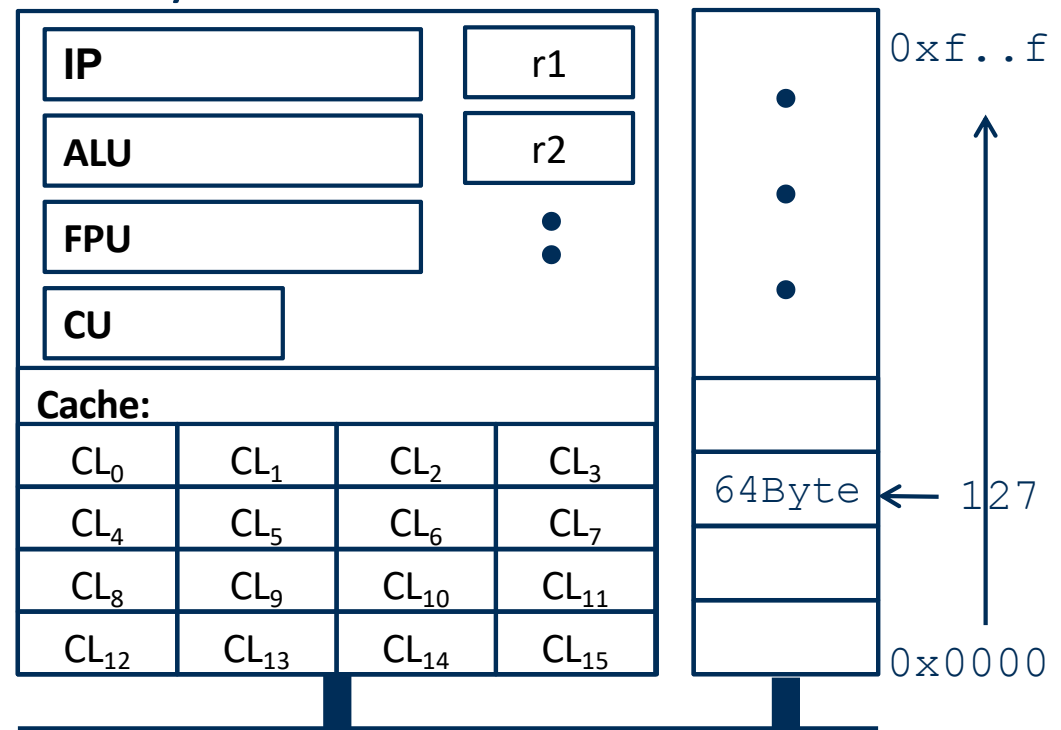
# CACHE: FUNCTION & STRUCTURE 1/2

Caches need to be fast!

For every memory access, CPU needs to see whether data is in Cache (Cache-Hit) or needs to fetch from memory (Cache-Miss)

Availability is checked using the memory accesses' **address**.

- Organized in Cache-Lines (CL)  
e.g. of 64 Bytes length;  
With a size of 1kB, there  
would be 16 Cache-Lines!
- Upon read from memory address  
127, every CL would have to  
be checked: “fully associative”
- To be practical, CLs are  
associated with specific memory addresses (so called Associativity)  
See next slides for implementations

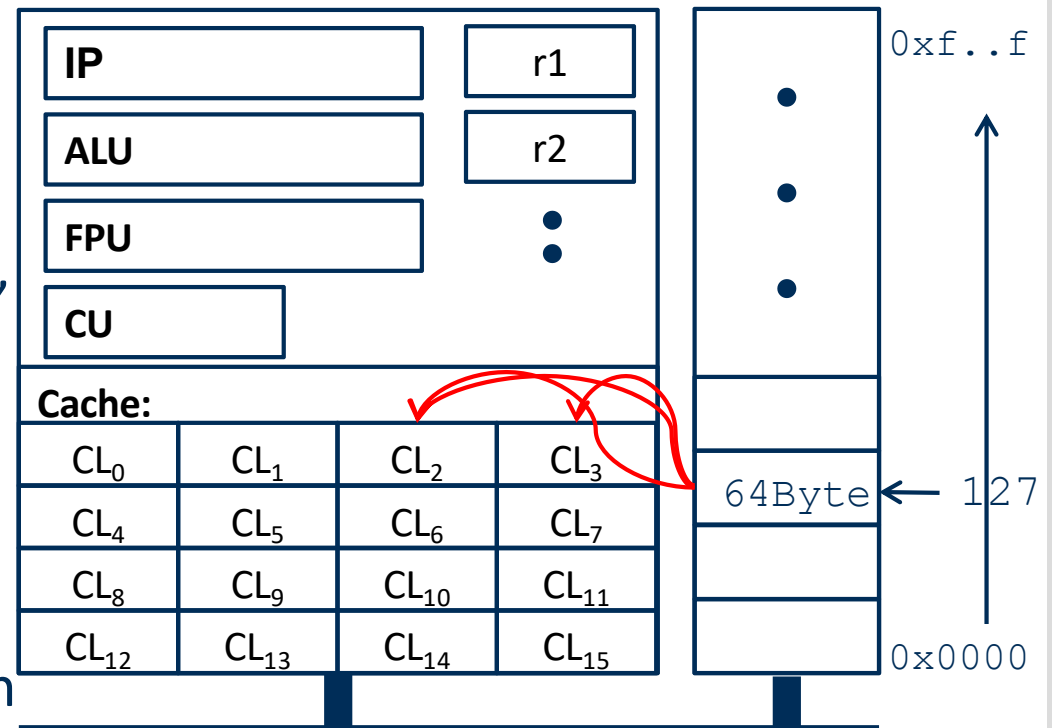


# CACHE: FUNCTION & STRUCTURE 2/2

Fully Associative (any CL may take any of the memory's blocks) are unpractical due to high power consumption to check all locations.

“One-way” associativity / Direct mapping : A memory block has only **one** possibly CL where it may be stored; not good if we have more than 16 blocks of memory...

- Two-way associativity:  
a memory block may be in one of two possibly CLs, i.e. we need two address comparators, and some eviction strategy!
- Four-way: 4 possible locations for cache lines
- ...
- Fully associative **any** CL location

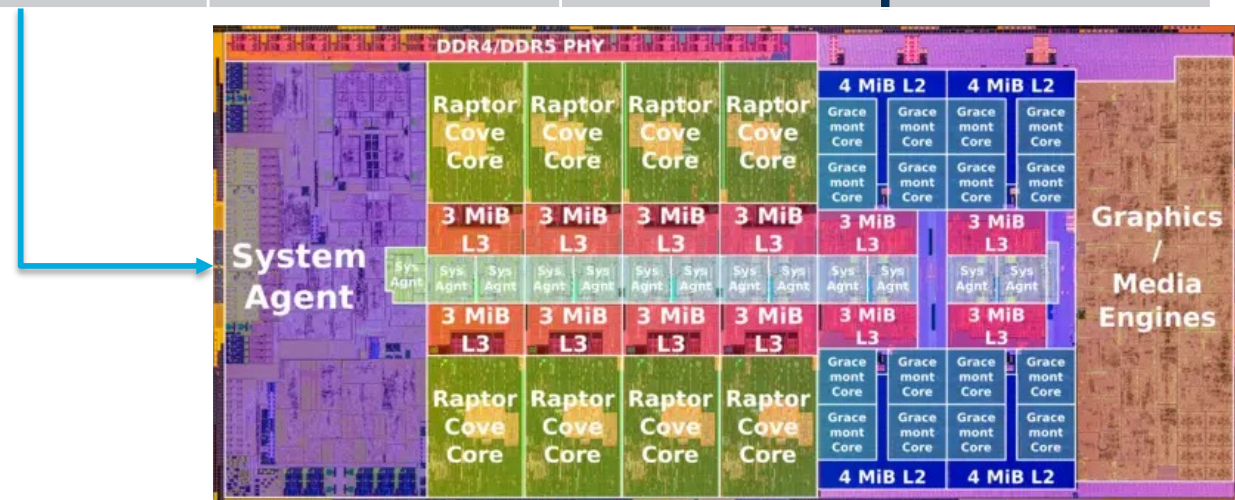




# CACHES: IMPLEMENTATIONS

CPU-local Caches are a key-component to achieve performance:  
The table lists generations of microprocessors (and it's cores/arch name)

Cache / Introduced?	Raptor Lake Core: Golden Cove 10/2022	Ice Lake Core: Sunny Cove 09/2019	Skylake Core: Broadwell 08/2015	AMD ZEN4 09/2022
L1 Data	48 kB, (12-way?)	48 kB, <b>12-way</b>	32 kB, 8-way	32 kB, 8-way
L1 Instruction	32 kB, (8-way?)	32 kB, 8-way	32 kB, 8-way	32 kB, 8-way
L2	<b>2 MB, 8-way (P)</b>	<b>512 kB, 8-way</b>	256 kB, 4-way	<b>1 MiB, 8-way</b>
L3/core	<b>3 MB, 16-way (P)</b>	2 MB, 16-way	2 MB, 16-way	<b>4 MB, 16-way</b>
μOp Cache	<b>4096</b>	2304	1536	<b>4096</b>



## CACHES: SIMPLE MODEL OF HOW MUCH TO GAIN

Assumption: Cache-Memory is  $\tau$  faster than main memory.


- In this simplistic model, we do not account for latency or bandwidth, but just look at time:  
$$T_m = \tau T_c$$
- With a Cache-Hit/Miss ratio  $\beta$ , i.e.  $\beta$  is the ratio of accesses to already cached data, what gain is to be expected?

- The average time is:

$$T_{av} = \beta T_c + (1 - \beta) T_m$$

- The gain  $G$  for variable  $\tau$  and  $\beta$  is therefore (shortened):

$$G(\tau, \beta) = \frac{T_m}{T_{av}} = \frac{\tau}{\beta + \tau(1 - \beta)}$$




## CACHES: CURRENT IMPLEMENTATIONS...

Caches are tailored for specific uses:

- Instruction Cache
- Data Cache

Both have specific access patterns and typical working set sizes.

E.g. the working set size of Integer portion of SPEC2000 Benchmark:



Source: Wikipedia

- Current implementations have up to 4 levels of Cache...
- AMD Ryzen Zen 3 with 3D Stacking „V-Cache“:
  - 64x 32kB Level 1 Instruction Cache (8-way associative), write-back
  - 64x 32kB Level 1 Data Cache (8-way associative), write-back
  - 64x 512kB Level 2 combined Cache (8-way associative), write-back
  - 64 to **256 MB** Level 3 combined Cache (16-way associative, with hash)
- Discuss, how Instruction Cache access patterns differ from Data

# SIMD


## SINGLE INSTRUCTION MULTIPLE DATA

Another big improvement in execution efficiency is SIMD:

1 instruction computes 2 (Intel SSE) or up to 64 (Intel AVX512) elements

This is an “old” concept: Vectorization (think NEC SX from 1990 ‘til now), compilers & programmers nevertheless need to adapt their code!

```
for (int i = 0; i < N; i++) {
    c[i] = a[i] + b[i];
}
```



Compiler needs to detect/allow:

- Loop unrolling (here by 4)  
(Is N divisible by 4? Handle N%4)
- Is N large enough?
- Are a, b, c properly aligned?  
Are they Pointers, and alias??

Result:

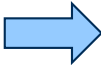
- ¼ of the (loop body’s) instructions (but not loop control).
- Better scheduling of core’s resources

# SIMD PROGRAMMING 1/2

Simple Loops, Simple data structures help!

## 1. Split Loops into multiple loops:


```
for (int i = 0; i < N; i++) {  
    if (a[i]%2)  
        c[i] = a[i] + b[i];  
    else  
        c[i] = a[i] - b[i];  
}
```



```
for (int i = 0; i < N; i++)  
    if (a[i]%2 != 0)  
        c[i] = a[i] + b[i];  
for (int i = 0; i < N; i++)  
    if (a[i]%2 == 0)  
        c[i] = a[i] - b[i];
```

## 2. Use Struct of Arrays (SoA) instead of Array of Structs (AoS):

```
struct {  
    int a;    // 4 Byte Padding  
    double c; // 8 Byte Aligned  
    float b;  
} aos[1000];
```



```
struct {  
    int a[1000];  
    float b[1000];  
    double c[1000];  
} soa;
```



# SIMD PROGRAMMING 2/2

Use the compiler to help & document the code:

Intel Compiler:

```
icc -vec-report=5 -c file.c    # Reports on vectorization of loops
icc -opt-report=5 -c file.c    # Reports optimization hints
```

GNU Compiler:

```
gcc -fverbose-asm -S file.c    # Generates verbose ASM output
gcc -fopt-info-vec-missed=vec.miss # Reports missed vectorization
```

Further options to get to vectorize:

1. Properly align types:

```
__declspec(align(64)) float a[1000];           // Windows
__attribute__((aligned(64))) float a[1000];    // Linux
and _mm_malloc() / _mm_free() for dynamic allocation:
float * c = _mm_malloc(1000*sizeof(float), 64);
```

2. Use (compiler-specific) intrinsics, e.g.

```
#include <immintrin.h>
_mm512_add_epi64 (__m512i __A, __m512i __B);
```

## SIMD RESOURCES & LINKS

The CPU vendors programming manual is a good starting point:

- Links for ARM's NEON and SVE/SVE2  
SVE: Scalable Vector Extension:
  - 32 Vector regs z0 ... z31, implementation defined: multiples of 128 Bit registers (with a max. of 2048 bits);
  - 16 predicate registers p0 ... p15 to select vector lanes.
- Intel C++ Compiler Guide on Vectorization
- Intel Vector Intrinsics Guide
- To check different in compiler capabilities, use the Compiler Explorer

# AMD APUs

## Am Beispiel AMD MI300A

## AMD APU


AMD Accelerated Processing Unit (APU) combine CPU & iGPU **since 2012** (e.g. used in PlayStation 4 & 5, Steam Deck)

- The main advantage is shared addressing space
- Due to the different design (ISA, memory access patterns and memory synchronization) it's a Heterogeneous System Architecture (HSA)

(news June / November 2023):

@SC2023: AMD Instinct MI300A (APUs)  
power LLNL's HPC "EL Capitan":

- **2 ExaFlops**, i.e.  $2 \cdot 10^{18}$  FP32 OP/s
- Due to massive Power of 24 core Zen 4 &
- GPU: CDNA-3 &
- Massive Memory Bandwidth (to 128 GB HBM3 RAM on package): **5,218 TB/s!!!**
- **According to AMD ML benchmarks 8x faster than MI250 GPUs**



## LECTURE: LITERATURE

According to list of literature:

- Patterson, D., Hennessy, J.: *Computer Organization and Design*, Kaufmann, 2011  
(Deutsche Übersetzung: Rechnerorganisation und –entwurf, Spektrum)
- Hennessy, J., Patterson, D.: *Computer Architecture: A quantitative Approach*, **5<sup>th</sup> edition**, Morgan Kaufmann, 2017
- Tanenbaum, A., Austin, T.: *Structured Computer Organization*, Pearson, 6<sup>th</sup> edition, 2013
- Tanenbaum, A., Austin, T.: *Rechnerarchitektur: von der digitalen Logik zum Parallelrechner*, Pearson, 6<sup>th</sup> ed., 2014
- Huang, H.W.: *The HCS12/9S12. An Introduction to the HW and SW interface*, Thomson Learning, 2009
- Beierlein, T.: *Taschenbuch Mikroprozessortechnik*, Carl-Hanser Verl., 2011

